

```
In [21]: # Ryan Picariello - 822856548 - Intro to ML Homework 1 part 1b
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [22]: df = pd.read_csv('C:/Users/Ryanj/Downloads/Housing.csv')
df.head() # To get first n rows from the dataset default value of n is 5
M=len(df)
```

```
In [23]: housing = pd.DataFrame(pd.read_csv('C:/Users/Ryanj/Downloads/Housing.csv'))
housing.head()
```

```
Out[23]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating
0	13300000	7420	4	2	3	yes	no	no	no
1	12250000	8960	4	4	4	yes	no	no	no
2	12250000	9960	3	2	2	yes	no	yes	no
3	12215000	7500	4	2	2	yes	no	yes	no
4	11410000	7420	4	1	2	yes	yes	yes	no

```
In [24]: # You can see that your dataset has many columns with values as 'Yes' or 'No'.
# But in order to fit a regression line, we would need numerical values and not string.
# List of variables to map
varlist = ['mainroad', 'guestroom', 'basement', 'hotwaterheating', 'airconditioning', '']
# Defining the map function
def binary_map(x):
    return x.map({'yes': 1, "no": 0})
# Applying the function to the housing list
housing[varlist] = housing[varlist].apply(binary_map)
# Check the housing dataframe now
housing.head()
```

```
Out[24]:
```

	price	area	bedrooms	bathrooms	stories	mainroad	guestroom	basement	hotwaterheating
0	13300000	7420	4	2	3	1	0	0	0
1	12250000	8960	4	4	4	1	0	0	0
2	12250000	9960	3	2	2	1	0	1	0
3	12215000	7500	4	2	2	1	0	1	0
4	11410000	7420	4	1	2	1	1	1	0

```
In [25]: #Splitting the Data into Training and Testing Sets
from sklearn.model_selection import train_test_split
# We specify this so that the train and test data set always have the same rows, respec
```

```
np.random.seed(0)
df_train, df_test = train_test_split(housing, train_size = 0.7, test_size = 0.3, random
```

```
In [26]: num_vars = ['area', 'bedrooms', 'bathrooms', 'mainroad', 'guestroom', 'basement', 'hotw
df_Newtrain = df_train[num_vars]
df_Newtest = df_test[num_vars]
df_Newtrain.head()
```

```
Out[26]:
```

	area	bedrooms	bathrooms	mainroad	guestroom	basement	hotwaterheating	airconditioning
454	4500	3	1	1	0	0	0	1
392	3990	3	1	1	0	0	0	0
231	4320	3	1	1	0	0	0	0
271	1905	5	1	0	0	1	0	0
250	3510	3	1	1	0	0	0	0

```
In [27]: XTrain = df_Newtrain.values[:,0:10]
YTrain = df_Newtrain.values[:,10]

XTest = df_Newtest.values[:,0:10]
YTest = df_Newtest.values[:,10]
```

```
In [28]: mean = np.ones(XTrain.shape[1])
std = np.ones(XTrain.shape[1])
for i in range(0, XTrain.shape[1]):
    mean[i] = np.mean(XTrain.transpose()[i])
    std[i] = np.std(XTrain.transpose()[i])
    for j in range(0, XTrain.shape[0]):
        XTrain[j][i] = (XTrain[j][i] - mean[i])/std[i]
```

```
In [29]: mean = np.ones(XTest.shape[1])
std = np.ones(XTest.shape[1])
for i in range(0, XTest.shape[1]):
    mean[i] = np.mean(XTest.transpose()[i])
    std[i] = np.std(XTest.transpose()[i])
    for j in range(0, XTest.shape[0]):
        XTest[j][i] = (XTest[j][i] - mean[i])/std[i]
```

```
In [30]: def compute_cost(X, n, theta):
    h = np.ones((X.shape[0],1))
    theta = theta.reshape(1,n+1)
    for i in range(0,X.shape[0]):
        h[i] = float(np.matmul(theta, X[i]))
    h = h.reshape(X.shape[0])
    return h
```

```
In [31]: def gradient_descent(X, y, theta, alpha, iterations, n, h):
    cost = np.ones(iterations)
```

```

for i in range(0, iterations):
    theta[0] = theta[0] - (alpha/X.shape[0]) * sum(h - y)
    for j in range(1, n+1):
        theta[j] = theta[j] - (alpha/X.shape[0]) * sum((h-y) * X.transpose()[j])
    h = compute_cost(X, n, theta)
    cost[i] = (1/X.shape[0]) * 0.5 * sum(np.square(h - y))
theta = theta.reshape(1, n+1)
return theta, cost

```

```

In [32]: def linear_regression(X, y, alpha, iterations):
        n = X.shape[1]
        one_column = np.ones((X.shape[0], 1))
        X = np.concatenate((one_column, X), axis = 1)
        theta = np.zeros(n+1)
        h = compute_cost(X, n, theta)
        theta, cost = gradient_descent(X, y, theta, alpha, iterations, n, h)
        return theta, cost

```

```

In [33]: iterations = 500;
        alpha = 0.1;
        alpha2 = 0.01

```

```

In [34]: ThetaTraining, CostTraining = linear_regression(XTrain, YTrain, alpha, iterations)
        print('Final value of theta with an alpha of ', alpha, ' =', ThetaTraining)
        CostTraining = list(CostTraining)
        nIterations_Training = [x for x in range(1, (iterations + 1))]

```

```

Final value of theta with an alpha of 0.1 = [[3670731.06244055  551896.57649369  38579
5.47482823  992491.61099674
 406487.69378465  296569.19055127  109169.67597818  340894.17909666
1248570.78075159  254848.24455347  899005.83201632]]

```

```

In [35]: ThetaTraining2, CostTraining2 = linear_regression(XTrain, YTrain, alpha2, iterations)
        print('Final value of theta with an alpha of ', alpha2, ' =', ThetaTraining2)
        CostTraining2 = list(CostTraining2)
        nIterations_Training2 = [x for x in range(1, (iterations + 1))]

```

```

Final value of theta with an alpha of 0.01 = [[3237989.41217222  547127.46438364  2470
32.84559761  1022333.00298265
 153039.07459016  364712.32134447  470973.32733094  381655.9909572
1307691.88614664  604186.7049498  911391.41717596]]

```

```

In [36]: theta_Test, cost_Test = linear_regression(XTest, YTest, alpha, iterations)
        print('Final value of theta with an alpha of ', alpha, ' =', theta_Test)
        cost_Test = list(cost_Test)
        nIterations_Test = [x for x in range(1, (iterations + 1))]

```

```

Final value of theta with an alpha of 0.1 = [[3665283.80998753  741402.85068115  26465
7.01674763  992948.00356364
 305680.32344454  126509.594599  265573.12956302  96887.61766261
1318395.63714537  574221.3020385  498783.75949985]]

```

```

In [37]: theta_Test2, cost_Test2 = linear_regression(XTest, YTest, alpha2, iterations)
        print('Final value of theta with an alpha of ', alpha2, ' =', theta_Test2)


```

```
cost_Test2 = list(cost_Test2)
nIterations_Test2 = [x for x in range(1,(iterations + 1))]
```

Final value of theta with an alpha of 0.01 = [[3211733.75281949 791345.02851107 162838.95231366 1164613.49088194
51779.18076014 239993.96046932 566736.95476624 137642.82043416
1204854.90843431 783960.095531 689075.93949973]]

```
In [39]: plt.plot(nIterations_Test, cost_Test, label='Test set Alpha 0.1')
plt.plot(nIterations_Test2, cost_Test2, label='Test set Alpha 0.01')
plt.plot(nIterations_Training, CostTraining, label='Traing set Alpha 0.1')
plt.plot(nIterations_Training2, CostTraining2, label='Traing set Alpha 0.01')
plt.legend()
plt.rcParams["figure.figsize"]=(10,6)
plt.grid()
plt.xlabel('No. of iterations')
plt.ylabel('Cost')
plt.title('Convergence of gradient descent')
```

Out[39]: Text(0.5, 1.0, 'Convergence of gradient descent')



In []: